

# Report

## Problem 1

1. Please print the model architecture of method A and B.

- difference
  - model: DCGAN -> SN-GAN spectral\_normalization (nn.utils.spectral\_norm)
  - optimizer: Adam -> AdamW
  - label smoothing
  - image noise
  - learning rate
- model A
  - generator

```
DCGAN_Generator(  
  (main): Sequential(  
    (0): ConvTranspose2d(100, 512, kernel_size=(4, 4), stride=(1,  
1), bias=False)  
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,  
track_running_stats=True)  
    (2): ReLU(inplace=True)  
    (3): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2,  
2), padding=(1, 1), bias=False)  
    (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,  
track_running_stats=True)  
    (5): ReLU(inplace=True)  
    (6): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2,  
2), padding=(1, 1), bias=False)  
    (7): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,  
track_running_stats=True)  
    (8): ReLU(inplace=True)  
    (9): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2),  
padding=(1, 1), bias=False)  
    (10): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,  
track_running_stats=True)  
    (11): ReLU(inplace=True)  
    (12): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2),  
padding=(1, 1), bias=False)  
    (13): Tanh()  
  )  
)
```

Layer (type)	Output Shape	Param #
ConvTranspose2d-1	[-1, 512, 4, 4]	819,200
BatchNorm2d-2	[-1, 512, 4, 4]	1,024
ReLU-3	[-1, 512, 4, 4]	0
ConvTranspose2d-4	[-1, 256, 8, 8]	2,097,152
BatchNorm2d-5	[-1, 256, 8, 8]	512

ReLU-6	[-1, 256, 8, 8]	0
ConvTranspose2d-7	[-1, 128, 16, 16]	524,288
BatchNorm2d-8	[-1, 128, 16, 16]	256
ReLU-9	[-1, 128, 16, 16]	0
ConvTranspose2d-10	[-1, 64, 32, 32]	131,072
BatchNorm2d-11	[-1, 64, 32, 32]	128
ReLU-12	[-1, 64, 32, 32]	0
ConvTranspose2d-13	[-1, 3, 64, 64]	3,072
Tanh-14	[-1, 3, 64, 64]	0
=====		
Total params: 3,576,704		
Trainable params: 3,576,704		
Non-trainable params: 0		
-----		
Input size (MB): 0.00		
Forward/backward pass size (MB): 3.00		
Params size (MB): 13.64		
Estimated Total Size (MB): 16.64		
-----		

#### ■ discriminator

DCGAN_Discriminator( (feature_extract): Sequential( (0): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False) (1): LeakyReLU(negative_slope=0.2, inplace=True) (2): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False) (3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True) (4): LeakyReLU(negative_slope=0.2, inplace=True) (5): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False) (6): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True) (7): LeakyReLU(negative_slope=0.2, inplace=True) (8): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False) (9): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True) (10): LeakyReLU(negative_slope=0.2, inplace=True) (11): Conv2d(512, 1, kernel_size=(4, 4), stride=(1, 1), bias=False) (12): sigmoid() ) ) )		
-----		
Layer (type)	Output shape	Param #
=====		
Conv2d-1	[-1, 64, 32, 32]	3,072
LeakyReLU-2	[-1, 64, 32, 32]	0
Conv2d-3	[-1, 128, 16, 16]	131,072
BatchNorm2d-4	[-1, 128, 16, 16]	256
LeakyReLU-5	[-1, 128, 16, 16]	0

Conv2d-6	[-1, 256, 8, 8]	524,288
BatchNorm2d-7	[-1, 256, 8, 8]	512
LeakyReLU-8	[-1, 256, 8, 8]	0
Conv2d-9	[-1, 512, 4, 4]	2,097,152
BatchNorm2d-10	[-1, 512, 4, 4]	1,024
LeakyReLU-11	[-1, 512, 4, 4]	0
Conv2d-12	[-1, 1, 1, 1]	8,192
Sigmoid-13	[-1, 1, 1, 1]	0
=====		
Total params: 2,765,568		
Trainable params: 2,765,568		
Non-trainable params: 0		
-----		
Input size (MB): 0.05		
Forward/backward pass size (MB): 2.31		
Params size (MB): 10.55		
Estimated Total Size (MB): 12.91		
-----		

- o model B

- generator

```
My_Generator(
  (main): Sequential(
    (0): ConvTranspose2d(100, 512, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace=True)
    (3): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (5): ReLU(inplace=True)
    (6): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (7): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (8): ReLU(inplace=True)
    (9): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (10): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (11): ReLU(inplace=True)
    (12): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (13): Tanh()
  )
)
```

-----		
Layer (type)	Output shape	Param #
=====		
ConvTranspose2d-1	[-1, 512, 4, 4]	819,200

BatchNorm2d-2	[-1, 512, 4, 4]	1,024
ReLU-3	[-1, 512, 4, 4]	0
ConvTranspose2d-4	[-1, 256, 8, 8]	2,097,152
BatchNorm2d-5	[-1, 256, 8, 8]	512
ReLU-6	[-1, 256, 8, 8]	0
ConvTranspose2d-7	[-1, 128, 16, 16]	524,288
BatchNorm2d-8	[-1, 128, 16, 16]	256
ReLU-9	[-1, 128, 16, 16]	0
ConvTranspose2d-10	[-1, 64, 32, 32]	131,072
BatchNorm2d-11	[-1, 64, 32, 32]	128
ReLU-12	[-1, 64, 32, 32]	0
ConvTranspose2d-13	[-1, 3, 64, 64]	3,072
Tanh-14	[-1, 3, 64, 64]	0
=====		
Total params: 3,576,704		
Trainable params: 3,576,704		
Non-trainable params: 0		
-----		
Input size (MB): 0.00		
Forward/backward pass size (MB): 3.00		
Params size (MB): 13.64		
Estimated Total Size (MB): 16.64		
-----		

#### ■ discriminator

```

My_Discriminator(
  (feature_extract): Sequential(
    (0): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=
(1, 1), bias=False)
    (1): LeakyReLU(negative_slope=0.2, inplace=True)
    (2): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=
(1, 1), bias=False)
    (3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (4): LeakyReLU(negative_slope=0.2, inplace=True)
    (5): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2),
padding=(1, 1), bias=False)
    (6): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (7): LeakyReLU(negative_slope=0.2, inplace=True)
    (8): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2),
padding=(1, 1), bias=False)
    (9): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (10): LeakyReLU(negative_slope=0.2, inplace=True)
    (11): Conv2d(512, 1, kernel_size=(4, 4), stride=(1, 1),
bias=False)
    (12): Sigmoid()
  )
)

```

Layer (type)	Output Shape	Param #
=====		
Conv2d-1	[-1, 64, 32, 32]	3,072

LeakyReLU-2	[-1, 64, 32, 32]	0
Conv2d-3	[-1, 128, 16, 16]	131,072
BatchNorm2d-4	[-1, 128, 16, 16]	256
LeakyReLU-5	[-1, 128, 16, 16]	0
Conv2d-6	[-1, 256, 8, 8]	524,288
BatchNorm2d-7	[-1, 256, 8, 8]	512
LeakyReLU-8	[-1, 256, 8, 8]	0
Conv2d-9	[-1, 512, 4, 4]	2,097,152
BatchNorm2d-10	[-1, 512, 4, 4]	1,024
LeakyReLU-11	[-1, 512, 4, 4]	0
Conv2d-12	[-1, 1, 1, 1]	8,192
Sigmoid-13	[-1, 1, 1, 1]	0
=====		
Total params: 2,765,568		
Trainable params: 2,765,568		
Non-trainable params: 0		
-----		
Input size (MB): 0.05		
Forward/backward pass size (MB): 2.31		
Params size (MB): 10.55		
Estimated Total Size (MB): 12.91		
-----		

2. Please show the first 32 generated images of both method A and B then discuss the difference between method A and B.

o A



- B



- difference
  - A整體比較白
  - B整體臉的占比比較大
  - A的四周白色或黑色區塊比較多也比較大
  - A的第二排比較多崩壞照，B沒什麼崩壞照

3. Please discuss what you've observed and learned from implementing GAN.

- model 越深不一定越好
- GAN不是和把圖片變太多
- 網路上說D可以train比較多step，learning rate也適合比較大，但似乎是通常是D太弱，而本次作業感覺是D偏強，所以那些training tricks反而比較沒有幫助。
- $D(x)$  和  $D(G(z))$ 越接近0.5越好，但太接近0.5就不太會繼續train了，而 $D(x)$ 太接近0.9也不太會繼續train
- 偶爾把real label和fake label交換對GAN的幫助不顯著

## Problem 2

1. Please print your model architecture and describe your implementation details.

- Unet

```
UNet(
  (time_embedding): TimeEmbedding(
    (time_embedding): Sequential(
      (0): Embedding(200, 128)
      (1): Linear(in_features=128, out_features=512, bias=True)
      (2): Swish()
      (3): Linear(in_features=512, out_features=512, bias=True)
    )
  )
  (cond_embedding): ConditionalEmbedding(
    (conditional_embedding): Sequential(
      (0): Embedding(11, 128, padding_idx=0)
      (1): Linear(in_features=128, out_features=512, bias=True)
      (2): Swish()
      (3): Linear(in_features=512, out_features=512, bias=True)
    )
  )
)
```

```

(head): Conv2d(3, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
(downblocks): ModuleList(
  (0): ResidualBlock(
    (block1): Sequential(
      (0): GroupNorm(32, 128, eps=1e-05, affine=True)
      (1): Swish()
      (2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (time_embedding_proj): Sequential(
      (0): Swish()
      (1): Linear(in_features=512, out_features=128, bias=True)
    )
    (cond_proj): Sequential(
      (0): Swish()
      (1): Linear(in_features=512, out_features=128, bias=True)
    )
    (block2): Sequential(
      (0): GroupNorm(32, 128, eps=1e-05, affine=True)
      (1): Swish()
      (2): Dropout(p=0.15, inplace=False)
      (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (shortcut): Identity()
    (attention): AttentionBlock(
      (group_norm): GroupNorm(32, 128, eps=1e-05, affine=True)
      (proj_q): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1))
      (proj_k): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1))
      (proj_v): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1))
      (proj): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1))
    )
  )
  (1): ResidualBlock(
    (block1): Sequential(
      (0): GroupNorm(32, 128, eps=1e-05, affine=True)
      (1): Swish()
      (2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (time_embedding_proj): Sequential(
      (0): Swish()
      (1): Linear(in_features=512, out_features=128, bias=True)
    )
    (cond_proj): Sequential(
      (0): Swish()
      (1): Linear(in_features=512, out_features=128, bias=True)
    )
    (block2): Sequential(
      (0): GroupNorm(32, 128, eps=1e-05, affine=True)
      (1): Swish()
      (2): Dropout(p=0.15, inplace=False)
      (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))

```

```

    )
    (shortcut): Identity()
    (attention): AttentionBlock(
      (group_norm): GroupNorm(32, 128, eps=1e-05, affine=True)
      (proj_q): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1))
      (proj_k): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1))
      (proj_v): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1))
      (proj): Conv2d(128, 128, kernel_size=(1, 1), stride=(1, 1))
    )
  )
  (2): DownSample(
    (conv_1): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1))
    (conv_2): Conv2d(128, 128, kernel_size=(5, 5), stride=(2, 2),
padding=(2, 2))
  )
  (3): ResidualBlock(
    (block1): Sequential(
      (0): GroupNorm(32, 128, eps=1e-05, affine=True)
      (1): Swish()
      (2): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (time_embedding_proj): Sequential(
      (0): Swish()
      (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (cond_proj): Sequential(
      (0): Swish()
      (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (block2): Sequential(
      (0): GroupNorm(32, 256, eps=1e-05, affine=True)
      (1): Swish()
      (2): Dropout(p=0.15, inplace=False)
      (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (shortcut): Conv2d(128, 256, kernel_size=(1, 1), stride=(1, 1))
    (attention): AttentionBlock(
      (group_norm): GroupNorm(32, 256, eps=1e-05, affine=True)
      (proj_q): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
      (proj_k): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
      (proj_v): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
      (proj): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
    )
  )
  (4): ResidualBlock(
    (block1): Sequential(
      (0): GroupNorm(32, 256, eps=1e-05, affine=True)
      (1): Swish()
      (2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (time_embedding_proj): Sequential(

```



```

        (0): Swish()
        (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (cond_proj): Sequential(
        (0): Swish()
        (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (block2): Sequential(
        (0): GroupNorm(32, 256, eps=1e-05, affine=True)
        (1): Swish()
        (2): Dropout(p=0.15, inplace=False)
        (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (shortcut): Identity()
    (attention): AttentionBlock(
        (group_norm): GroupNorm(32, 256, eps=1e-05, affine=True)
        (proj_q): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
        (proj_k): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
        (proj_v): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
        (proj): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
    )
    )
    (5): DownSample(
        (conv_1): Conv2d(256, 256, kernel_size=(3, 3), stride=(2, 2),
padding=(1, 1))
        (conv_2): Conv2d(256, 256, kernel_size=(5, 5), stride=(2, 2),
padding=(2, 2))
    )
    (6): ResidualBlock(
        (block1): Sequential(
            (0): GroupNorm(32, 256, eps=1e-05, affine=True)
            (1): Swish()
            (2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
        )
        (time_embedding_proj): Sequential(
            (0): Swish()
            (1): Linear(in_features=512, out_features=256, bias=True)
        )
        (cond_proj): Sequential(
            (0): Swish()
            (1): Linear(in_features=512, out_features=256, bias=True)
        )
        (block2): Sequential(
            (0): GroupNorm(32, 256, eps=1e-05, affine=True)
            (1): Swish()
            (2): Dropout(p=0.15, inplace=False)
            (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
        )
        (shortcut): Identity()
        (attention): AttentionBlock(
            (group_norm): GroupNorm(32, 256, eps=1e-05, affine=True)
            (proj_q): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))

```

```

        (proj_k): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
        (proj_v): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
        (proj): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
    )
)
(7): ResidualBlock(
  (block1): Sequential(
    (0): GroupNorm(32, 256, eps=1e-05, affine=True)
    (1): Swish()
    (2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
  )
  (time_embedding_proj): Sequential(
    (0): Swish()
    (1): Linear(in_features=512, out_features=256, bias=True)
  )
  (cond_proj): Sequential(
    (0): Swish()
    (1): Linear(in_features=512, out_features=256, bias=True)
  )
  (block2): Sequential(
    (0): GroupNorm(32, 256, eps=1e-05, affine=True)
    (1): Swish()
    (2): Dropout(p=0.15, inplace=False)
    (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
  )
  (shortcut): Identity()
  (attention): AttentionBlock(
    (group_norm): GroupNorm(32, 256, eps=1e-05, affine=True)
    (proj_q): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
    (proj_k): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
    (proj_v): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
    (proj): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
  )
)
)
(middleblocks): ModuleList(
  (0): ResidualBlock(
    (block1): Sequential(
      (0): GroupNorm(32, 256, eps=1e-05, affine=True)
      (1): Swish()
      (2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (time_embedding_proj): Sequential(
      (0): Swish()
      (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (cond_proj): Sequential(
      (0): Swish()
      (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (block2): Sequential(
      (0): GroupNorm(32, 256, eps=1e-05, affine=True)

```

```

        (1): Swish()
        (2): Dropout(p=0.15, inplace=False)
        (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (shortcut): Identity()
    (attention): AttentionBlock(
        (group_norm): GroupNorm(32, 256, eps=1e-05, affine=True)
        (proj_q): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
        (proj_k): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
        (proj_v): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
        (proj): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1))
    )
)
(1): ResidualBlock(
    (block1): Sequential(
        (0): GroupNorm(32, 256, eps=1e-05, affine=True)
        (1): Swish()
        (2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (time_embedding_proj): Sequential(
        (0): Swish()
        (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (cond_proj): Sequential(
        (0): Swish()
        (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (block2): Sequential(
        (0): GroupNorm(32, 256, eps=1e-05, affine=True)
        (1): Swish()
        (2): Dropout(p=0.15, inplace=False)
        (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (shortcut): Identity()
    (attention): Identity()
)
)
(upblocks): ModuleList(
  (0): ResidualBlock(
    (block1): Sequential(
      (0): GroupNorm(32, 512, eps=1e-05, affine=True)
      (1): Swish()
      (2): Conv2d(512, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (time_embedding_proj): Sequential(
      (0): Swish()
      (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (cond_proj): Sequential(
      (0): Swish()
      (1): Linear(in_features=512, out_features=256, bias=True)

```

```

    )
    (block2): Sequential(
      (0): GroupNorm(32, 256, eps=1e-05, affine=True)
      (1): Swish()
      (2): Dropout(p=0.15, inplace=False)
      (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (shortcut): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1))
    (attention): Identity()
  )
  (1): ResidualBlock(
    (block1): Sequential(
      (0): GroupNorm(32, 512, eps=1e-05, affine=True)
      (1): Swish()
      (2): Conv2d(512, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (time_embedding_proj): Sequential(
      (0): Swish()
      (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (cond_proj): Sequential(
      (0): Swish()
      (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (block2): Sequential(
      (0): GroupNorm(32, 256, eps=1e-05, affine=True)
      (1): Swish()
      (2): Dropout(p=0.15, inplace=False)
      (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (shortcut): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1))
    (attention): Identity()
  )
  (2): ResidualBlock(
    (block1): Sequential(
      (0): GroupNorm(32, 512, eps=1e-05, affine=True)
      (1): Swish()
      (2): Conv2d(512, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (time_embedding_proj): Sequential(
      (0): Swish()
      (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (cond_proj): Sequential(
      (0): Swish()
      (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (block2): Sequential(
      (0): GroupNorm(32, 256, eps=1e-05, affine=True)
      (1): Swish()
      (2): Dropout(p=0.15, inplace=False)

```

```

        (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (shortcut): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1))
    (attention): Identity()
)
(3): UpSample(
    (conv): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    (convtrans): ConvTranspose2d(256, 256, kernel_size=(5, 5), stride=
(2, 2), padding=(2, 2), output_padding=(1, 1))
)
(4): ResidualBlock(
    (block1): Sequential(
        (0): GroupNorm(32, 512, eps=1e-05, affine=True)
        (1): Swish()
        (2): Conv2d(512, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (time_embedding_proj): Sequential(
        (0): Swish()
        (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (cond_proj): Sequential(
        (0): Swish()
        (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (block2): Sequential(
        (0): GroupNorm(32, 256, eps=1e-05, affine=True)
        (1): Swish()
        (2): Dropout(p=0.15, inplace=False)
        (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (shortcut): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1))
    (attention): Identity()
)
(5): ResidualBlock(
    (block1): Sequential(
        (0): GroupNorm(32, 512, eps=1e-05, affine=True)
        (1): Swish()
        (2): Conv2d(512, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (time_embedding_proj): Sequential(
        (0): Swish()
        (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (cond_proj): Sequential(
        (0): Swish()
        (1): Linear(in_features=512, out_features=256, bias=True)
    )
    (block2): Sequential(
        (0): GroupNorm(32, 256, eps=1e-05, affine=True)
        (1): Swish()

```

```

        (2): Dropout(p=0.15, inplace=False)
        (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (shortcut): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1))
    (attention): Identity()
)
(6): ResidualBlock(
  (block1): Sequential(
    (0): GroupNorm(32, 384, eps=1e-05, affine=True)
    (1): Swish()
    (2): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
  )
  (time_embedding_proj): Sequential(
    (0): Swish()
    (1): Linear(in_features=512, out_features=256, bias=True)
  )
  (cond_proj): Sequential(
    (0): Swish()
    (1): Linear(in_features=512, out_features=256, bias=True)
  )
  (block2): Sequential(
    (0): GroupNorm(32, 256, eps=1e-05, affine=True)
    (1): Swish()
    (2): Dropout(p=0.15, inplace=False)
    (3): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
  )
  (shortcut): Conv2d(384, 256, kernel_size=(1, 1), stride=(1, 1))
  (attention): Identity()
)
(7): UpSample(
  (conv): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
  (convtrans): ConvTranspose2d(256, 256, kernel_size=(5, 5), stride=
(2, 2), padding=(2, 2), output_padding=(1, 1))
)
(8): ResidualBlock(
  (block1): Sequential(
    (0): GroupNorm(32, 384, eps=1e-05, affine=True)
    (1): Swish()
    (2): Conv2d(384, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
  )
  (time_embedding_proj): Sequential(
    (0): Swish()
    (1): Linear(in_features=512, out_features=128, bias=True)
  )
  (cond_proj): Sequential(
    (0): Swish()
    (1): Linear(in_features=512, out_features=128, bias=True)
  )
  (block2): Sequential(
    (0): GroupNorm(32, 128, eps=1e-05, affine=True)

```

```

        (1): Swish()
        (2): Dropout(p=0.15, inplace=False)
        (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
    )
    (shortcut): Conv2d(384, 128, kernel_size=(1, 1), stride=(1, 1))
    (attention): Identity()
)
(9): ResidualBlock(
  (block1): Sequential(
    (0): GroupNorm(32, 256, eps=1e-05, affine=True)
    (1): Swish()
    (2): Conv2d(256, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
  )
  (time_embedding_proj): Sequential(
    (0): Swish()
    (1): Linear(in_features=512, out_features=128, bias=True)
  )
  (cond_proj): Sequential(
    (0): Swish()
    (1): Linear(in_features=512, out_features=128, bias=True)
  )
  (block2): Sequential(
    (0): GroupNorm(32, 128, eps=1e-05, affine=True)
    (1): Swish()
    (2): Dropout(p=0.15, inplace=False)
    (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
  )
  (shortcut): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1))
  (attention): Identity()
)
(10): ResidualBlock(
  (block1): Sequential(
    (0): GroupNorm(32, 256, eps=1e-05, affine=True)
    (1): Swish()
    (2): Conv2d(256, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
  )
  (time_embedding_proj): Sequential(
    (0): Swish()
    (1): Linear(in_features=512, out_features=128, bias=True)
  )
  (cond_proj): Sequential(
    (0): Swish()
    (1): Linear(in_features=512, out_features=128, bias=True)
  )
  (block2): Sequential(
    (0): GroupNorm(32, 128, eps=1e-05, affine=True)
    (1): Swish()
    (2): Dropout(p=0.15, inplace=False)
    (3): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
  )
)

```

```


        (shortcut): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1))
        (attention): Identity()
    )
)
(tail): Sequential(
  (0): GroupNorm(32, 128, eps=1e-05, affine=True)
  (1): Swish()
  (2): Conv2d(128, 3, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1))
)
)

```

◦ train process

- 把time和label(condition)進行embedding
- 每個時間點都會加上隨機但亮在固定範圍的雜訊
- image forward Unet到最後會生成image size的latent
- latent和常態分布取loss · 希望學出常態分布
- 從latent生成圖片時會去看統計量去加noise

2. Please show 10 generated images for each digit (0-9) in your report. You can put all 100 outputs in one image with columns indicating different noise inputs and rows indicating different digits.



3. Visualize total six images in the reverse process of the first "0" in your grid in (2) with different time steps.

- T: 0 40 80 120 160 200



4. Please discuss what you've observed and learned from implementing conditional diffusion model.

- 因為想讓不同時間點同時訓練，tensor會擴增時間維度，時間越久所需空間也就越大
- 統計值是看各個時間點各種condition的整組，所以batch之中可能會互相影響到，也就代表batch size也變得重要
- 慢慢加noise是一種幫助generate的手段，而能不能train好主要跟model比較相關

## Problem 3

1. Please create and fill the table with the following format in your report




	MNIST-M $\rightarrow$ SVHN	MNIST-M $\rightarrow$ USPS
Trained on source	0.42852646818153206	0.719758064516129
Adaptation (DANN)	0.4484798892175993	0.7681451612903226
Trained on target	0.8558569899918173	0.8487903225806451

2. Please visualize the latent space of DANN by mapping the validation images to 2D space with t-SNE. For each scenario, you need to plot two figures which are colored by digit class (0-9) and by domain, respectively.


- Note that you need to plot the figures of both 2 scenarios, so 4 figures in total.

- SVHN

- by class




- by domain




- USPS

■ by class



■ by domain



3. Please describe the implementation details of your model and discuss what you've observed and learned from implementing DANN.

```
DANN_model(  
  (feature_extract): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(5, 5), stride=(1, 1))  
    (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,  
track_running_stats=True)  
    (2): ReLU(inplace=True)  
    (3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,  
ceil_mode=False)  
    (4): Conv2d(64, 48, kernel_size=(5, 5), stride=(1, 1))  
    (5): BatchNorm2d(48, eps=1e-05, momentum=0.1, affine=True,  
track_running_stats=True)  
    (6): Dropout2d(p=0.3, inplace=True)  
    (7): ReLU(inplace=True)
```

```

(8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
)
(class_classifier): Sequential(
  (0): Dropout(p=0.3, inplace=False)
  (1): Linear(in_features=768, out_features=192, bias=True)
  (2): BatchNorm1d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (3): Mish(inplace=True)
  (4): Dropout(p=0.3, inplace=False)
  (5): Linear(in_features=192, out_features=128, bias=True)
  (6): BatchNorm1d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (7): Mish(inplace=True)
  (8): Linear(in_features=128, out_features=10, bias=True)
)
(domain_classifier): Sequential(
  (0): Dropout(p=0.3, inplace=False)
  (1): Linear(in_features=768, out_features=128, bias=True)
  (2): BatchNorm1d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (3): Mish(inplace=True)
  (4): Linear(in_features=128, out_features=2, bias=True)
)
)

```

Layer (type)	Output Shape	Param #
=====		
Conv2d-1	[-1, 64, 24, 24]	4,864
BatchNorm2d-2	[-1, 64, 24, 24]	128
ReLU-3	[-1, 64, 24, 24]	0
MaxPool2d-4	[-1, 64, 12, 12]	0
Conv2d-5	[-1, 48, 8, 8]	76,848
BatchNorm2d-6	[-1, 48, 8, 8]	96
Dropout2d-7	[-1, 48, 8, 8]	0
ReLU-8	[-1, 48, 8, 8]	0
MaxPool2d-9	[-1, 48, 4, 4]	0
Dropout-10	[-1, 768]	0
Linear-11	[-1, 192]	147,648
BatchNorm1d-12	[-1, 192]	384
Mish-13	[-1, 192]	0
Dropout-14	[-1, 192]	0
Linear-15	[-1, 128]	24,704
BatchNorm1d-16	[-1, 128]	256
Mish-17	[-1, 128]	0
Linear-18	[-1, 10]	1,290
Dropout-19	[-1, 768]	0
Linear-20	[-1, 128]	98,432
BatchNorm1d-21	[-1, 128]	256
Mish-22	[-1, 128]	0
Linear-23	[-1, 2]	258
=====		

Total params: 355,164

Trainable params: 355,164

Non-trainable params: 0

```
-----  
Input size (MB): 0.01  
Forward/backward pass size (MB): 1.04  
Params size (MB): 1.35  
Estimated Total Size (MB): 2.40  
-----
```

- detail
  - 用convolution作為feature extraction，取的latent後，分別有class跟domain的classifier。訓練domain時要把gradient反轉，用alpha作為反轉的開關
  - alpha和現在的訓練進度有關，依據當前在總共要訓練的資料是中的比例來調整
  - model 主要還是在學class classifier，domain classifie會讓class classifier不會學到太好
  - 認為DANN幫助有限
- MNISTM -> SVHN
  - class\_classifier不宜太難
  - regularization大比較合適
  - 難train是因為task難很多而且差異domain很大(數字數量)，所以model太複雜容易over fitting
- MNISTM -> USPS
  - class\_classifier深比較好
  - 很快就overfitting